

HSplit interface

Inria TADaaM Team

19 septembre 2018

1 Communicator/group creation functions

1.1 MPIX_COMM_SPLIT_TYPE

`MPI_COMM_SPLIT_TYPE` is part of MPI 3.1 and is a communicator management routine defined in Chapter 6 of the standard (Groups, Contexts, Communicators and Caching), section 6.4 (Communicator Management).

```
MPI_COMM_SPLIT_TYPE(comm, split_type, key, info, newcomm)
```

IN `comm` communicator (handle)

IN `split_type` type of processes to be grouped together (integer)

IN `key` control of rank assignment (integer)

IN `info` info argument (handle)

OUT `newcomm` new communicator (handle)

C Prototype :

```
int MPI_Comm_split_type(MPI_Comm comm, int split_type, int key,
                        MPI_Info info, MPI_Comm *newcomm)
```

MPI 3.1 defines a single value for `split_type` : `MPI_COMM_TYPE_SHARED`, that allows the function to split the original communicator into subcommunicators, which encompass MPI processes that can create a shared memory region with the other member processes.

Current extension/implementation :

- Introduces/relies on a new `split_type` value : `MPI_COMM_TYPE_HW_DOMAIN` Behaviour description : the input communicator is split into subcommunicators, each of which encompasses MPI processes that do share resources in the underlying physical topology (e.g. a network switch, a physical node, a L3 cache, a L2 cache, a core, etc.). In order to avoid the creation of redundant objects, the group of MPI processes supporting a new subcommunicator should be a *strict* subset of the group supporting the input (parent) communicator. More specifically, a call to `MPI_COMM_COMPARE(comm, newcomm)` should return `MPI_UNEQUAL`. In case of several possible levels, the *lowest* level in the hierarchy should be considered and intermediate levels

skipped. If no valid communicator can be created, the value `MPI_COMM_NULL` is returned, and this value can be used to assess if the last level of the hierarchy has been reached. In particular, it is possible for an implementation to not produce subcommunicators by directly returning the value `MPI_COMM_NULL`.

It is possible to capture the hierarchical nature of the underlying hardware by calling "recursively" `MPI_COMM_SPLIT_TYPE` with `MPI_COMM_TYPE_HW_DOMAIN` as `split_type` value on newly created subcommunicators. For instance :

```
MPI_Comm newcomm[32];
MPI_Comm oldcomm = MPI_COMM_WORLD;
int rank, idx = 0;
while((oldcomm != MPI_COMM_NULL) && (idx < 32)){
    MPI_Comm_rank(oldcomm, &rank);
    MPI_Comm_split_type(oldcomm,
                        MPI_COMM_TYPE_HW_DOMAIN,
                        rank,
                        MPI_INFO_NULL,
                        &newcomm[idx]);
    oldcomm = newcomm[idx++];
}
```

The mapping and binding of MPI processes onto physical resources must be taken into account for subcommunicators creation. Indeed, the deepest hardware level corresponding to a subcommunicator should always correspond to the binding of the calling MPI process. For instance, if a process is bound to a certain cache level, no information below this cache level can be returned, as the MPI process can possibly use any of the caches below the level it is bound to. Any attempt to create a subcommunicator corresponding to a hardware level below the level the MPI process is bound to return `MPI_COMM_NULL`.

— Uses several *keyvals* :

— `MPI_HW_DOMAIN_TYPE` : If this key is defined (in the `info` parameter passed to the function), `MPI_COMM_SPLIT_TYPE` can then perform a split operation on a specific, desired hardware level, whose name is the value of the `MPI_HW_DOMAIN_TYPE` key. The levels names are not specified and are implementation-dependent. Such names can be queried by the `MPIX_GET_HW_DOMAIN_INFO` and `MPIX_GET_HW_TOPOLOGY_INFO` routines. For instance, an HWLOC-based implementation could use the HWLOC types names as domain types :

- `HWLOC_OBJ_MACHINE`
- `HWLOC_OBJ_PACKAGE` (former `HWLOC_OBJ_SOCKET`)
- `HWLOC_OBJ_CORE`
- `HWLOC_OBJ_PU` for hardware threads
- `HWLOC_OBJ_NUMANODE`
- `HWLOC_OBJ_L1CACHE`, ..., `HWLOC_OBJ_L5CACHE` (new terminology coming in future releases)
- `HWLOC_OBJ_GROUP` for other hierarchy levels

Simpler names (e.g. `CORE`, `SOCKET`, `Machine`) should be acceptable also. In this way, names are still not specified in the MPI standard but can be used nevertheless.

— `MPI_HW_DOMAIN_NUM` : This key is used internally and is never made available to the

user. It indicates the number of subcommunicators produced after a split operation at a specific level of the hardware hierarchy. This information can then be used to distribute data among the newly created subcommunicators. The value can be obtained by the user with a call to `MPIX_GET_HW_DOMAIN_INFO`.

- `MPI_HW_DOMAIN_RANK` : This key is used internally and is never made available to the user. It indicates the "rank" of a subcommunicator produced after a split operation at a specific level of the hardware hierarchy. This information can then be used to distribute data among the newly created subcommunicators. The value can be obtained by the user with a call to `MPIX_GET_HW_DOMAIN_INFO`.

1.2 MPPIX_GET_HW_DOMAIN_NEIGHBOURS

`MPIX_GET_HW_DOMAIN_NEIGHBOURS(comm,hops,metric,newgroup)`

IN `comm` communicator (handle)

IN `hops` number of hops in the topology, neighborhood extent (integer)

IN `metric` neighborhood type (integer)

OUT `newgroup` new group (handle)

C Prototype :

```
int MPPIX_Get_hw_domain_neighbours(MPI_Comm comm, int hops,
                                   int metric,MPI_Group *newgroup)
```

Current extension/implementation :

- Creates a *group* instead of a *communicator* since it is probably a local operation and the information should not propagate to other processes in the input communicator.
- Can possibly create neighborhood of various types, for instance *memory/computing* neighborhoods, *network* neighborhoods, etc.

2 Query functions

2.1 MPPIX_GET_HW_DOMAIN_INFO

`MPIX_GET_HW_DOMAIN_INFO(comm,num_subcomms,rank,type,info)`

IN `comm` communicator (handle)

OUT `num_subcomms` number of subcommunicators (integer)

OUT `rank` domain "rank" (integer)

OUT `type` domain name (string)

IN `info` info object (handle)

C Prototype :

```
int MPIX_Get_hw_domain_info(MPI_Comm comm,int *num_subcomms,
                           int *rank,char **type, MPI_Info info)
```

Current extension/implementation :

- Uses internally the `MPI_HW_DOMAIN_TYPE`, `MPI_HW_DOMAIN_NUM`, and `MPI_HW_DOMAIN_RANK` keys defined in `MPI_COMM_SPLIT_TYPE`, and returns their values to the user in a more usable form.
- Uses an `info` parameter that should not be in the final version of the interface, as this object should be attached to the input communicator with the `MPI_Comm_set_info` function in `MPI_COMM_SPLIT_TYPE` and retrieved with the `MPI_Comm_get_info` function in `MPIX_GET_HW_DOMAIN_INFO`.

2.2 MPIX_GET_HW_TOPOLOGY_INFO

```
MPIX_GET_HW_TOPOLOGY_INFO(numlevels,info)
```

OUT `numlevels` number of levels in the hardware hierarchy (integer)

OUT `info` info object (handle)

C Prototype :

```
int MPIX_Get_hw_topology_info(int *numlevels,MPI_Info info)
```

Current extension/implementation :

allows the calling MPI process to retrieve information about the underlying hardware topology. Two different types of information are available :

- the number of hardware levels in the hierarchy (the `numlevels` parameter)
- a set of keyvals (set in the `info` parameter). There are *numlevels* different keyvals, named `MPI_HW_LEVEL0`, `MPI_HW_LEVEL1`, ... `MPI_HW_LEVELnumlevels-1`. These values are implementation-defined and can be the same as the ones used to define the `MPI_HW_DOMAIN_TYPE` key(s) in `MPIX_COMM_SPLIT_TYPE`.

2.3 MPIX_GET_MIN_HW_DOMAIN

```
MPIX_GET_MIN_HW_DOMAIN(comm,size,ranks,type)
```

IN `comm` communicator (handle)

IN `size` size of ranks array (integer)

IN `ranks` ranks (array of integer)

OUT `type` hardware domain name (string)

C Prototype :

```
int MPIX_Get_min_hw_domain(MPI_Comm comm, int size, int *ranks, char **type)
```

Current extension/implementation :

Returns the name of the *lowest* level in the hierarchy shared by all the MPI processes which ranks in the communicator `comm` are listed in the `ranks` array. If the calling process rank is not among the ranks listed in the array passed as an argument, the type returned should be "Unknown" or "Invalid".

3 Code example

```
{
    MPI_Comm out_comm;
    MPI_Info info;
    int rank, scomm_num, scomm_rank;
    char *resource_type = NULL;

    MPI_Info_create(&info);
    MPI_Info_set(info, "MPI_HW_DOMAIN_TYPE", "Socket");
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPIX_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_HW_DOMAIN, rank, info, &out_comm);

    /* check info about subcomm */
    if (out_comm != MPI_COMM_NULL){
        MPIX_Get_hw_domain_info(out_comm, &scomm_num, &scomm_rank, &resource_type, info);
        fprintf(stdout, "=== Number of subcomms : %i\n", scomm_num);
        fprintf(stdout, "=== Subcomm rank : %i\n", scomm_rank);
        fprintf(stdout, "=== Subcomm type : %s\n", resource_type);
    } else {
        fprintf(stdout, "No level found\n");
    }
}
```