

*D R A F T*

# Document for a Standard Message-Passing Interface

Message Passing Interface Forum

October 7, 2021

This work was supported in part by NSF and ARPA under NSF contract CDA-9115428 and Esprit under project HPC Standards (21111).

This is the result of a LaTeX run of a draft of a single chapter of the MPIF Final Report document.

## Chapter 8

# MPI Environmental Management

This chapter discusses routines for getting and, where appropriate, setting various parameters that relate to the MPI implementation and the execution environment (such as error handling). The procedures for entering and leaving the MPI execution environment are also described here.

### 8.1 Implementation Information

#### 8.1.1 Version Inquiries

In order to cope with changes to the MPI Standard, there are both compile-time and run-time ways to determine which version of the standard is in use in the environment one is using.

The “version” will be represented by two separate integers, for the version and subversion: In C,

```
#define MPI_VERSION    4
#define MPI_SUBVERSION 0
```

in Fortran,

```
INTEGER :: MPI_VERSION, MPI_SUBVERSION
PARAMETER (MPI_VERSION    = 4)
PARAMETER (MPI_SUBVERSION = 0)
```

For runtime determination,

**MPI\_GET\_VERSION(version, subversion)**

OUT	version	version number (integer)
OUT	subversion	subversion number (integer)

#### C binding

```
int MPI_Get_version(int *version, int *subversion)
```

#### Fortran 2008 binding

```
MPI_Get_version(version, subversion, ierror)
```

```

1      INTEGER, INTENT(OUT) :: version, subversion
2      INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

### Fortran binding

```

4      MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)
5      INTEGER VERSION, SUBVERSION, IERROR

```

MPI\_GET\_VERSION can be called at any time in an MPI program. This function must always be thread-safe, as defined in Section 11.6. Valid (MPI\_VERSION, MPI\_SUBVERSION) pairs in this and previous versions of the MPI standard are (4,0), (3,1), (3,0), (2,2), (2,1), (2,0), and (1,2).

```

13     MPI_GET_LIBRARY_VERSION(version, resultlen)

```

14	OUT	version	version number (string)
15	OUT	resultlen	Length (in printable characters) of the result
16			returned in <b>version</b> (integer)

### C binding

```

19     int MPI_Get_library_version(char *version, int *resultlen)

```

### Fortran 2008 binding

```

22     MPI_Get_library_version(version, resultlen, ierror)
23     CHARACTER(LEN=MPI_MAX_LIBRARY_VERSION_STRING), INTENT(OUT) :: version
24     INTEGER, INTENT(OUT) :: resultlen
25     INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

### Fortran binding

```

28     MPI_GET_LIBRARY_VERSION(VERSION, RESULTLEN, IERROR)
29     CHARACTER*(*) VERSION
30     INTEGER RESULTLEN, IERROR

```

This routine returns a string representing the version of the MPI library. The version argument is a character string for maximum flexibility.

*Advice to implementors.* An implementation of MPI should return a different string for every change to its source code or build that could be visible to the user. (*End of advice to implementors.*)

The argument **version** must represent storage that is MPI\_MAX\_LIBRARY\_VERSION\_STRING characters long. MPI\_GET\_LIBRARY\_VERSION may write up to this many characters into **version**.

The number of characters actually written is returned in the output argument, **resultlen**. In C, a null character is additionally stored at **version[resultlen]**. The value of **resultlen** cannot be larger than MPI\_MAX\_LIBRARY\_VERSION\_STRING - 1. In Fortran, **version** is padded on the right with blank characters. The value of **resultlen** cannot be larger than MPI\_MAX\_LIBRARY\_VERSION\_STRING.

MPI\_GET\_LIBRARY\_VERSION can be called at any time in an MPI program. This function must always be thread-safe, as defined in Section 11.6.

### 8.1.2 Environmental Inquiries

When using the World Model (Section 11.2), a set of attributes that describe the execution environment is attached to the communicator `MPI_COMM_WORLD` when MPI is initialized. The values of these attributes can be inquired by using the function `MPI_COMM_GET_ATTR` described in Section 7.7 and in Section 19.3.7. It is erroneous to delete these attributes, free their keys, or change their values.

The list of predefined attribute keys include

**MPI\_TAG\_UB** Upper bound for tag value.

**MPI\_HOST** Host process rank, if such exists, `MPI_PROC_NULL`, otherwise.

**MPI\_IO** rank of a node that has regular I/O facilities (possibly myrank). Nodes in the same communicator may return different values for this parameter.

**MPI\_WTIME\_IS\_GLOBAL** Boolean variable that indicates whether clocks are synchronized.

When using the Sessions Model (Section 11.3), only the `MPI_TAG_UB` attribute is available.

Vendors may add implementation-specific parameters (such as node number, real memory size, virtual memory size, etc.)

These predefined attributes do not change value between MPI initialization (`MPI_INIT`) and MPI completion (`MPI_FINALIZE`), and cannot be updated or deleted by users.

*Advice to users.* Note that in the C binding, the value returned by these attributes is a *pointer* to an `int` containing the requested value. (*End of advice to users.*)

The required parameter values are discussed in more detail below:

#### Tag Values

Tag values range from 0 to the value returned for `MPI_TAG_UB`, inclusive. These values are guaranteed to be unchanging during the execution of an MPI program. In addition, the tag upper bound value must be *at least* 32767. An MPI implementation is free to make the value of `MPI_TAG_UB` larger than this; for example, the value  $2^{30} - 1$  is also a valid value for `MPI_TAG_UB`.

In the Sessions Model, the attribute `MPI_TAG_UB` is attached to all communicators created by `MPI_COMM_CREATE_FROM_GROUP` and `MPI_INTERCOMM_CREATE_FROM_GROUPS`, with the same value on all MPI processes in the communicator. In the World Model, the attribute `MPI_TAG_UB` has the same value on all processes of `MPI_COMM_WORLD`.

#### Host Rank

The value returned for `MPI_HOST` gets the rank of the *HOST* process in the group associated with communicator `MPI_COMM_WORLD`, if there is such. `MPI_PROC_NULL` is returned if there is no host. MPI does not specify what it means for a process to be a *HOST*, nor does it require that a *HOST* exists.

The attribute `MPI_HOST` has the same value on all processes of `MPI_COMM_WORLD`.

## IO Rank

The value returned for `MPI_IO` is the rank of a processor that can provide language-standard I/O facilities. For Fortran, this means that all of the Fortran I/O operations are supported (e.g., `OPEN`, `REWIND`, `WRITE`). For C, this means that all of the ISO C I/O operations are supported (e.g., `fopen`, `fprintf`, `lseek`).

If every process can provide language-standard I/O, then the value `MPI_ANY_SOURCE` will be returned. Otherwise, if the calling process can provide language-standard I/O, then its rank will be returned. Otherwise, if some process can provide language-standard I/O then the rank of one such process will be returned. The same value need not be returned by all processes. If no process can provide language-standard I/O, then the value `MPI_PROC_NULL` will be returned.

*Advice to users.* Note that input is not collective, and this attribute does *not* indicate which process can or does provide input. (*End of advice to users.*)

## Clock Synchronization

The value returned for `MPI_WTIME_IS_GLOBAL` is 1 if clocks at all processes in `MPI_COMM_WORLD` are synchronized, 0 otherwise. A collection of clocks is considered synchronized if explicit effort has been taken to synchronize them. The expectation is that the variation in time, as measured by calls to `MPI_WTIME`, will be less than one half the round-trip time for an MPI message of length zero. If time is measured at a process just before a send and at another process just after a matching receive, the second time should be always higher than the first one.

The attribute `MPI_WTIME_IS_GLOBAL` need not be present when the clocks are not synchronized (however, the attribute key `MPI_WTIME_IS_GLOBAL` is always valid). This attribute may be associated with communicators other than `MPI_COMM_WORLD`.

The attribute `MPI_WTIME_IS_GLOBAL` has the same value on all processes of `MPI_COMM_WORLD`.

## Inquire Processor Name

`MPI_GET_PROCESSOR_NAME(name, resultlen)`

OUT	name	A unique specifier for the actual (as opposed to virtual) node.
OUT	resultlen	Length (in printable characters) of the result returned in name

## C binding

```
int MPI_Get_processor_name(char *name, int *resultlen)
```

## Fortran 2008 binding

```
MPI_Get_processor_name(name, resultlen, ierror)
  CHARACTER(LEN=MPI_MAX_PROCESSOR_NAME), INTENT(OUT) :: name
  INTEGER, INTENT(OUT) :: resultlen
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

**Fortran binding**

```

MPI_GET_PROCESSOR_NAME(NAME, RESULTLEN, IERROR)
    CHARACTER*(*) NAME
    INTEGER RESULTLEN, IERROR

```

This routine returns the name of the processor on which it was called at the moment of the call. The name is a character string for maximum flexibility. From this value it must be possible to identify a specific piece of hardware; possible values include “processor 9 in rack 4 of mpp.cs.org” and “231” (where 231 is the actual processor number in the running homogeneous system). The argument `name` must represent storage that is at least `MPI_MAX_PROCESSOR_NAME` characters long. `MPI_GET_PROCESSOR_NAME` may write up to this many characters into `name`.

The number of characters actually written is returned in the output argument, `resultlen`. In C, a null character is additionally stored at `name[resultlen]`. The value of `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME-1`. In Fortran, `name` is padded on the right with blank characters. The value of `resultlen` cannot be larger than `MPI_MAX_PROCESSOR_NAME`.

*Rationale.* This function allows MPI implementations that do process migration to return the current processor. Note that nothing in MPI *requires* or defines process migration; this definition of `MPI_GET_PROCESSOR_NAME` simply allows such an implementation. (*End of rationale.*)

*Advice to users.* The user must provide at least `MPI_MAX_PROCESSOR_NAME` space to write the processor name—processor names can be this long. The user should examine the output argument, `resultlen`, to determine the actual length of the name. (*End of advice to users.*)

**Inquire Hardware Resource Types and Status**

```

MPI_GET_HW_RESOURCE_TYPES(hw_info)

```

```

    OUT      hw_info          info object created (handle)

```

**C binding**

```

int MPI_Get_hw_resource_types(MPI_Info *hw_info)

```

**Fortran 2008 binding**

```

MPI_Get_hw_resource_types(hw_info, ierror)
    TYPE(MPI_Info), INTENT(OUT) :: hw_info
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

**Fortran binding**

```

MPI_GET_HW_RESOURCE_TYPES(HW_INFO, IERROR)
    INTEGER HW_INFO, IERROR

```

`MPI_GET_HW_RESOURCE_TYPES` returns an info object containing information pertaining to the hardware platform on which the calling MPI process is executing at the moment of the call. The information is stored in the following info keys:

- "mpi\_hw\_res\_nresources" is an integer that represents the number of all hardware resource types recognized by the MPI implementation and to which the calling MPI process can be restricted.
- "mpi\_hw\_res\_i\_type" is the type of the  $i$ -th hardware resource to which the calling MPI process can be restricted (with  $i \in \{0, \dots, \text{"mpi\_hw\_res\_nresources"} - 1\}$ ).
- "mpi\_hw\_res\_i\_aliases" is an integer that represents the number of hardware resource types that are aliases to "mpi\_hw\_res\_i\_type" (with  $i \in \{0, \dots, \text{"mpi\_hw\_res\_nresources"} - 1\}$ ).
- "mpi\_hw\_res\_i\_alias\_k" with  $k \in \{0, \dots, \text{"mpi\_hw\_res\_i\_aliases"} - 1\}$  is an integer  $j$  (with  $j \in \{0, \dots, \text{"mpi\_hw\_res\_nresources"} - 1\}$ ) such that "mpi\_hw\_res\_j\_type" is an alias to "mpi\_hw\_res\_i\_type".
- "mpi\_hw\_res\_i\_occupied", where  $i \in \{0, \dots, \text{"mpi\_hw\_res\_nresources"} - 1\}$ , is "true" if the calling MPI process is restricted to an instance of a hardware resource of type "mpi\_hw\_res\_i\_type" at the moment of the call. Otherwise, it is "false".

This local routine can return different information if the MPI processes are executing on different hardware resources. The user is responsible for freeing hw\_info via MPI\_INFO\_FREE.

*Advice to users.* The types returned by this routine can be used in MPI\_COMM\_SPLIT\_TYPE with the split\_type value MPI\_COMM\_TYPE\_HW\_GUIDED as key values for the info key "mpi\_hw\_resource\_type". However, the information returned in hw\_info may not be constant throughout the execution of the program because an MPI process can relocate (e.g., migrate or change its hardware restrictions). *(End of advice to users.)*

The following routine can be used to detect whether an MPI process is restricted to a hardware resource which type is returned by MPI\_GET\_HW\_RESOURCE\_TYPES.

**MPI\_GET\_HW\_RESOURCE\_STATUS(name, status)**

IN	name	name of hardware resource type (string)
OUT	status	status of the hardware resource type queried upon (integer)

#### C binding

```
int MPI_Get_hw_resource_status(char *name, int *status)
```

#### Fortran 2008 binding

```
MPI_Get_hw_resource_status(name, status, ierror)
  CHARACTER(LEN=MPI_MAX_INFO_KEY), INTENT(IN) :: name
  INTEGER, INTENT(OUT) :: status
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

#### Fortran binding

```
MPI_GET_HW_RESOURCE_STATUS(NAME, STATUS, IERROR)
```



```

CHARACTER*(*) NAME
INTEGER STATUS, IERROR

```

MPI\_HW\_PRESENT results if a hardware resource of type `name` is recognized but the calling MPI process cannot be restricted to an instance of a resource of this type (e.g., a network interface card). MPI\_HW\_USABLE results if the calling MPI process is able to be restricted to an instance of a hardware resource of type `name` at the moment of the call. MPI\_HW\_OCCUPIED results if the calling MPI process is actually restricted to an instance of a hardware resource of type `name` at the moment of the call. MPI\_HW\_UNKNOWN results otherwise.

If the status of a hardware resource is MPI\_HW\_OCCUPIED, its type can then be used as the value for the info key "mpi\_hw\_resource\_type" in the function MPI\_COMM\_SPLIT\_TYPE with the `split_type` value MPI\_COMM\_TYPE\_HW\_GUIDED. In that case, the splitting operation is guaranteed to produce a valid communicator for that MPI process.

## 8.2 Memory Allocation

In some systems, message-passing and remote-memory-access (RMA) operations run faster when accessing specially allocated memory (e.g., memory that is shared by the other processes in the communicating group on an SMP). MPI provides a mechanism for allocating and freeing such special memory. The use of such memory for message-passing or RMA is not mandatory, and this memory can be used without restrictions as any other dynamically allocated memory. However, implementations may restrict the use of some RMA functionality as defined in Section 12.5.3.

MPI\_ALLOC\_MEM(size, info, baseptr)

IN	size	size of memory segment in bytes (non-negative integer)
IN	info	info argument (handle)
OUT	baseptr	pointer to beginning of memory segment allocated

### C binding

```
int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)
```

### Fortran 2008 binding

```

MPI_Alloc_mem(size, info, baseptr, ierror)
  USE, INTRINSIC :: ISO_C_BINDING, ONLY : C_PTR
  INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: size
  TYPE(MPI_Info), INTENT(IN) :: info
  TYPE(C_PTR), INTENT(OUT) :: baseptr
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

### Fortran binding

```

MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)
  INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR
  INTEGER INFO, IERROR

```

If the Fortran compiler provides `TYPE(C_PTR)`, then the following generic interface must be provided in the `mpi` module and should be provided in `mpif.h` through overloading, i.e., with the same routine name as the routine with `INTEGER(KIND=MPI_ADDRESS_KIND)` `BASEPTR`, but with a different specific procedure name:

```

INTERFACE MPI_ALLOC_MEM
  SUBROUTINE MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)
    IMPORT :: MPI_ADDRESS_KIND
    INTEGER :: INFO, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) :: SIZE, BASEPTR
  END SUBROUTINE
  SUBROUTINE MPI_ALLOC_MEM_CPTR(SIZE, INFO, BASEPTR, IERROR)
    USE, INTRINSIC :: ISO_C_BINDING, ONLY : C_PTR
    IMPORT :: MPI_ADDRESS_KIND
    INTEGER :: INFO, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) :: SIZE
    TYPE(C_PTR) :: BASEPTR
  END SUBROUTINE
END INTERFACE

```

The base procedure name of this overloaded function is `MPI_ALLOC_MEM_CPTR`. The implied specific procedure names are described in Section 19.1.5.

By default, the allocated memory shall be aligned to at least the alignment required for load/store accesses of any datatype corresponding to a predefined MPI datatype. The `info` argument may be used to specify a desired alternative minimum alignment in bytes for the allocated memory by setting the value of the key `"mpi_minimum_memory_alignment"` to an integral number equal to a power of two. An implementation may ignore values smaller than the default required alignment. The `info` argument can also be used to provide directives that control the desired location of the allocated memory. Such a directive does not affect the semantics of the call. The corresponding `info` values are implementation-dependent. A null directive value of `info = MPI_INFO_NULL` is always valid.

The function `MPI_ALLOC_MEM` may return an error code of class `MPI_ERR_NO_MEM` to indicate it failed because memory is exhausted.

**MPI\_FREE\_MEM(base)**

<b>IN</b> <b>base</b>	initial address of memory segment allocated by MPI_ALLOC_MEM (choice)
-----------------------	--

#### C binding

```
int MPI_Free_mem(void *base)
```

#### Fortran 2008 binding

```

MPI_Free_mem(base, ierror)
  TYPE(*), DIMENSION(..), INTENT(IN), ASYNCHRONOUS :: base
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

#### Fortran binding

```
MPI_FREE_MEM(BASE, IERROR)
```

```
<type> BASE(*)
INTEGER IERROR
```

The function `MPI_FREE_MEM` may return an error code of class `MPI_ERR_BASE` to indicate an invalid base argument.

*Rationale.* The C bindings of `MPI_ALLOC_MEM` and `MPI_FREE_MEM` are similar to the bindings for the `malloc` and `free` C library calls: a call to `MPI_Alloc_mem(..., &base)` should be paired with a call to `MPI_Free_mem(base)` (one less level of indirection). Both arguments are declared to be of same type `void*` so as to facilitate type casting. The Fortran binding is consistent with the C bindings: the Fortran `MPI_ALLOC_MEM` call returns in `baseptr` the `TYPE(C_PTR)` pointer or the (integer valued) address of the allocated memory. The `base` argument of `MPI_FREE_MEM` is a choice argument, which passes (a reference to) the variable stored at that location. (*End of rationale.*)

*Advice to implementors.* If `MPI_ALLOC_MEM` allocates special memory, then a design similar to the design of C `malloc` and `free` functions has to be used, in order to find out the size of a memory segment, when the segment is freed. If no special memory is used, `MPI_ALLOC_MEM` simply invokes `malloc`, and `MPI_FREE_MEM` invokes `free`.

A call to `MPI_ALLOC_MEM` can be used in shared memory systems to allocate memory in a shared memory segment. (*End of advice to implementors.*)

**Example 8.1** Example of use of `MPI_ALLOC_MEM`, in Fortran with `TYPE(C_PTR)` pointers. We assume 4-byte REALs.

```
USE mpi_f08 ! or USE mpi      (not guaranteed with INCLUDE 'mpif.h')
USE, INTRINSIC :: ISO_C_BINDING
TYPE(C_PTR) :: p
REAL, DIMENSION(:,:), POINTER :: a ! no memory is allocated
INTEGER, DIMENSION(2) :: shape
INTEGER(KIND=MPI_ADDRESS_KIND) :: size
shape = (/100,100/)
size = 4 * shape(1) * shape(2) ! assuming 4 bytes per REAL
CALL MPI_Alloc_mem(size,MPI_INFO_NULL,p,ierr) ! memory is allocated and
CALL C_F_POINTER(p, a, shape) ! intrinsic ! now accessible via a(i,j)
... ! in ISO_C_BINDING
a(3,5) = 2.71
...
CALL MPI_Free_mem(a, ierr) ! memory is freed
```

**Example 8.2** Example of use of `MPI_ALLOC_MEM`, in Fortran with nonstandard *Cray-pointers*. We assume 4-byte REALs, and assume that these pointers are address-sized.

```
REAL A
POINTER (P, A(100,100)) ! no memory is allocated
INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
```

```

1  SIZE = 4*100*100
2  CALL MPI_ALLOC_MEM(SIZE, MPI_INFO_NULL, P, IERR)
3  ! memory is allocated
4  ...
5  A(3,5) = 2.71
6  ...
7  CALL MPI_FREE_MEM(A, IERR) ! memory is freed
8

```

This code is not Fortran 77 or Fortran 90 code. Some compilers may not support this code or need a special option, e.g., the GNU gFortran compiler needs `-fcray-pointer`.

*Advice to implementors.* Some compilers map Cray-pointers to address-sized integers, some to `TYPE(C_PTR)` pointers (e.g., Cray Fortran, version 7.3.3). From the user's viewpoint, this mapping is irrelevant because Examples 8.2 should work correctly with an MPI-3.0 (or later) library if Cray-pointers are available. (*End of advice to implementors.*)

**Example 8.3** Same example, in C.

```

19
20  float  (* f)[100][100];
21  /* no memory is allocated */
22  MPI_Alloc_mem(sizeof(float)*100*100, MPI_INFO_NULL, &f);
23  /* memory allocated */
24  ...
25  (*f)[5][3] = 2.71;
26  ...
27  MPI_Free_mem(f);
28

```

### 8.3 Error Handling

An MPI implementation may be unable or choose not to handle some failures that occur during MPI calls. These can include failures that generate exceptions or traps, such as floating point errors or access violations. The set of failures that are handled by MPI is implementation-dependent. Each such failure causes an error to be raised.

The above text takes precedence over any text on error handling within this document. Specifically, text that states that errors *will* be handled should be read as *may* be handled. More background information about how MPI treats errors can be found in Section 2.8.

A user can associate error handlers to four types of objects: communicators, windows, files, and sessions. The specified error handling routine will be used for any error that occurs during a call to MPI for the respective object. MPI calls that are not related to any MPI objects are considered to be attached to the communicator `MPI_COMM_SELF` when using the World Model (see Section 11.2). When `MPI_COMM_SELF` is not initialized (i.e., before `MPI_INIT` / `MPI_INIT_THREAD`, after `MPI_FINALIZE`, or when using the Sessions Model exclusively) the error raises the initial error handler (set during the launch operation, see 11.8.4). The attachment of error handlers to objects is purely local: different processes may attach different error handlers to corresponding objects.

Several predefined error handlers are available in MPI:

**MPI\_ERRORS\_ARE\_FATAL** The handler, when called, causes the program to abort all connected MPI processes. This is similar to calling `MPI_ABORT` using a communicator containing all connected processes with an implementation-specific value as the `errorcode` argument.

**MPI\_ERRORS\_ABORT** The handler, when called, is invoked on a communicator in a manner similar to calling `MPI_ABORT` on that communicator. If the error handler is invoked on an window or file, it is similar to calling `MPI_ABORT` using a communicator containing the group of MPI processes associated with the window or file, respectively. If the error handler is invoked on a session, the operation aborts only the local MPI process. In all cases, the value that would be provided as the `errorcode` argument to `MPI_ABORT` is implementation-specific.

**MPI\_ERRORS\_RETURN** The handler has no effect other than returning the error code to the user.

*Advice to implementors.* The implementation-specific error information resulting from `MPI_ERRORS_ARE_FATAL` and `MPI_ERRORS_ABORT` provided to the invoking environment should be meaningful to the end-user, for example a predefined error class. (*End of advice to implementors.*)

Implementations may provide additional predefined error handlers and programmers can code their own error handlers.

Unless otherwise requested, the error handler `MPI_ERRORS_ARE_FATAL` is set as the default initial error handler and associated with predefined communicators. Thus, if the user chooses not to control error handling, every error that MPI handles is treated as fatal. Since (almost) all MPI calls return an error code, a user may choose to handle errors in its main code, by testing the return code of MPI calls and executing a suitable recovery code when the call was not successful. In this case, the error handler `MPI_ERRORS_RETURN` will be used. Usually it is more convenient and more efficient not to test for errors after each MPI call, and have such error handled by a nontrivial MPI error handler. Note that unlike predefined communicators, windows and files do not inherit from the initial error handler, as defined in Sections 12.6 and 14.7 respectively.

When an error is raised, MPI will provide the user information about that error using an error code. Some errors might prevent MPI from completing further API calls successfully and those functions will continue to report errors until the cause of the error is corrected or the user terminates the application. The user can make the determination of whether or not to attempt to continue when handling such an error.

*Advice to users.* For example, users may be unable to correct errors corresponding to some error classes, such as `MPI_ERR_INTERRN`. Such errors may cause subsequent MPI calls to complete in error. (*End of advice to users.*)

*Advice to implementors.* A high-quality implementation will, to the greatest possible extent, circumscribe the impact of an error, so that normal processing can continue after an error handler was invoked. The implementation documentation will provide information on the possible effect of each class of errors and available recovery actions. (*End of advice to implementors.*)

An MPI error handler is an opaque object, which is accessed by a handle. MPI calls are provided to create new error handlers, to associate error handlers with objects, and to test which error handler is associated with an object. C has distinct typedefs for user defined error handling callback functions that accept communicator, file, window, and session arguments. In Fortran there are four user routines.

An error handler object is created by a call to `MPI_XXX_CREATE_ERRHANDLER`, where `XXX` is, respectively, `COMM`, `WIN`, `FILE`, or `SESSION`.

An error handler is attached to a communicator, window, file, or session by a call to `MPI_XXX_SET_ERRHANDLER`. The error handler must be either a predefined error handler, or an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`, with matching `XXX`. An error handler can also be attached to a session using the `errorhandler` argument to `MPI_SESSION_INIT`. The predefined error handlers `MPI_ERRORS_RETURN` and `MPI_ERRORS_ARE_FATAL` can be attached to communicators, windows, files, or sessions.

The error handler currently associated with a communicator, window, file, or session can be retrieved by a call to `MPI_XXX_GET_ERRHANDLER`.

The MPI function `MPI_ERRHANDLER_FREE` can be used to free an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`.

`MPI_XXX_GET_ERRHANDLER` behave as if a new error handler object is created. That is, once the error handler is no longer needed, `MPI_ERRHANDLER_FREE` should be called with the error handler returned from `MPI_XXX_GET_ERRHANDLER` to mark the error handler for deallocation. This provides behavior similar to that of `MPI_COMM_GROUP` and `MPI_GROUP_FREE`.

*Advice to implementors.* High-quality implementations should raise an error when an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER` is attached to an object of the wrong type with a call to `MPI_YYY_SET_ERRHANDLER`. To do so, it is necessary to maintain, with each error handler, information on the typedef of the associated user function. (*End of advice to implementors.*)

The syntax for these calls is given below.

### 8.3.1 Error Handlers for Communicators

`MPI_COMM_CREATE_ERRHANDLER(comm_errhandler_fn, errhandler)`

IN	<code>comm_errhandler_fn</code>	user defined error handling procedure (function)
OUT	<code>errhandler</code>	MPI error handler (handle)

#### C binding

```
int MPI_Comm_create_errhandler(
    MPI_Comm_errhandler_function *comm_errhandler_fn,
    MPI_Errhandler *errhandler)
```

#### Fortran 2008 binding

```
MPI_Comm_create_errhandler(comm_errhandler_fn, errhandler, ierror)
  PROCEDURE(MPI_Comm_errhandler_function) :: comm_errhandler_fn
  TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_COMM_CREATE_ERRHANDLER(COMM_ERRHANDLER_FN, ERRHANDLER, IERROR)
```

```
EXTERNAL COMM_ERRHANDLER_FN
```

```
INTEGER ERRHANDLER, IERROR
```

Creates an error handler that can be attached to communicators.

The user routine should be, in C, a function of type `MPI_Comm_errhandler_function`, which is defined as

```
typedef void MPI_Comm_errhandler_function(MPI_Comm *comm, int *error_code,
                                           ...);
```

The first argument is the communicator in use. The second is the error code to be returned by the MPI routine that raised the error. If the routine would have returned `MPI_ERR_IN_STATUS`, it is the error code returned in the status for the request that caused the error handler to be invoked. The remaining arguments are “`varargs`” arguments whose number and meaning is implementation-dependent. An implementation should clearly document these arguments. Addresses are used so that the handler may be written in Fortran. With the Fortran `mpi_f08` module, the user routine `comm_errhandler_fn` should be of the form:

```
ABSTRACT INTERFACE
```

```
  SUBROUTINE MPI_Comm_errhandler_function(comm, error_code)
```

```
    TYPE(MPI_Comm) :: comm
```

```
    INTEGER :: error_code
```

With the Fortran `mpi` module and `mpif.h`, the user routine `COMM_ERRHANDLER_FN` should be of the form:

```
SUBROUTINE COMM_ERRHANDLER_FUNCTION(COMM, ERROR_CODE)
```

```
  INTEGER COMM, ERROR_CODE
```

*Rationale.* The variable argument list is provided because it provides an ISO-standard hook for providing additional information to the error handler; without this hook, ISO C prohibits additional arguments. (*End of rationale.*)

*Advice to users.* A newly created communicator inherits the error handler that is associated with the “parent” communicator. In particular, the user can specify a “global” error handler for all communicators by associating this handler with the communicator `MPI_COMM_WORLD` immediately after initialization. (*End of advice to users.*)

```
MPI_COMM_SET_ERRHANDLER(comm, errhandler)
```

```
INOUT  comm                communicator (handle)
```

```
IN      errhandler          new error handler for communicator (handle)
```

### C binding

```
int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)
```

**Fortran 2008 binding**

```

MPI_Comm_set_errhandler(comm, errhandler, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

**Fortran binding**

```

MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR

```

Attaches a new error handler to a communicator. The error handler must be either a predefined error handler, or an error handler created by a call to MPI\_COMM\_CREATE\_ERRHANDLER.

```

MPI_COMM_GET_ERRHANDLER(comm, errhandler)

```

IN	comm	communicator (handle)
OUT	errhandler	error handler currently associated with communicator (handle)

**C binding**

```

int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)

```

**Fortran 2008 binding**

```

MPI_Comm_get_errhandler(comm, errhandler, ierror)
    TYPE(MPI_Comm), INTENT(IN) :: comm
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

**Fortran binding**

```

MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
    INTEGER COMM, ERRHANDLER, IERROR

```

Retrieves the error handler currently associated with a communicator.

For example, a library function may register at its entry point the current error handler for a communicator, set its own private error handler for this communicator, and restore before exiting the previous error handler.

**8.3.2 Error Handlers for Windows**

```

MPI_WIN_CREATE_ERRHANDLER(win_errhandler_fn, errhandler)

```

IN	win_errhandler_fn	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

**C binding**

```

int MPI_Win_create_errhandler(
    MPI_Win_errhandler_function *win_errhandler_fn,

```



```
MPI_Errhandler *errhandler)
```

### Fortran 2008 binding

```
MPI_Win_create_errhandler(win_errhandler_fn, errhandler, ierror)
  PROCEDURE(MPI_Win_errhandler_function) :: win_errhandler_fn
  TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_WIN_CREATE_ERRHANDLER(WIN_ERRHANDLER_FN, ERRHANDLER, IERROR)
  EXTERNAL WIN_ERRHANDLER_FN
  INTEGER ERRHANDLER, IERROR
```

Creates an error handler that can be attached to a window object. The user routine should be, in C, a function of type `MPI_Win_errhandler_function` which is defined as

```
typedef void MPI_Win_errhandler_function(MPI_Win *win, int *error_code,
    ...);
```

The first argument is the window in use, the second is the error code to be returned. The remaining arguments are “varargs” arguments whose number and meaning is implementation-dependent. An implementation should clearly document these arguments. With the Fortran `mpi_f08` module, the user routine `win_errhandler_fn` should be of the form:

```
ABSTRACT INTERFACE
  SUBROUTINE MPI_Win_errhandler_function(win, error_code)
    TYPE(MPI_Win) :: win
    INTEGER :: error_code
```

With the Fortran `mpi` module and `mpif.h`, the user routine `WIN_ERRHANDLER_FN` should be of the form:

```
SUBROUTINE WIN_ERRHANDLER_FUNCTION(WIN, ERROR_CODE)
  INTEGER WIN, ERROR_CODE
```

```
MPI_WIN_SET_ERRHANDLER(win, errhandler)
```

INOUT	win	window object (handle)
IN	errhandler	new error handler for window (handle)

### C binding

```
int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)
```

### Fortran 2008 binding

```
MPI_Win_set_errhandler(win, errhandler, ierror)
  TYPE(MPI_Win), INTENT(IN) :: win
  TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
  INTEGER WIN, ERRHANDLER, IERROR
```

Attaches a new error handler to a window. The error handler must be either a pre-defined error handler, or an error handler created by a call to `MPI_WIN_CREATE_ERRHANDLER`.

`MPI_WIN_GET_ERRHANDLER(win, errhandler)`

IN	win	window object (handle)
OUT	errhandler	error handler currently associated with window (handle)

### C binding

```
int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)
```

### Fortran 2008 binding

```
MPI_Win_get_errhandler(win, errhandler, ierror)
  TYPE(MPI_Win), INTENT(IN) :: win
  TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
  INTEGER WIN, ERRHANDLER, IERROR
```

Retrieves the error handler currently associated with a window.

## 8.3.3 Error Handlers for Files

`MPI_FILE_CREATE_ERRHANDLER(file_errhandler_fn, errhandler)`

IN	file_errhandler_fn	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

### C binding

```
int MPI_File_create_errhandler(
    MPI_File_errhandler_function *file_errhandler_fn,
    MPI_Errhandler *errhandler)
```

### Fortran 2008 binding

```
MPI_File_create_errhandler(file_errhandler_fn, errhandler, ierror)
  PROCEDURE(MPI_File_errhandler_function) :: file_errhandler_fn
  TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_FILE_CREATE_ERRHANDLER(FILE_ERRHANDLER_FN, ERRHANDLER, IERROR)
  EXTERNAL FILE_ERRHANDLER_FN
  INTEGER ERRHANDLER, IERROR
```

Creates an error handler that can be attached to a file object. The user routine should be, in C, a function of type `MPI_File_errhandler_function`, which is defined as

```
typedef void MPI_File_errhandler_function(MPI_File *file, int *error_code,
                                          ...);
```

The first argument is the file in use, the second is the error code to be returned. The remaining arguments are “`varargs`” arguments whose number and meaning is implementation-dependent. An implementation should clearly document these arguments.

With the Fortran `mpi_f08` module, the user routine `file_errhandler_fn` should be of the form:

ABSTRACT INTERFACE

```
SUBROUTINE MPI_File_errhandler_function(file, error_code)
```

```
  TYPE(MPI_File) :: file
```

```
  INTEGER :: error_code
```

With the Fortran `mpi` module and `mpif.h`, the user routine `FILE_ERRHANDLER_FN` should be of the form:

```
SUBROUTINE FILE_ERRHANDLER_FUNCTION(FILE, ERROR_CODE)
```

```
  INTEGER FILE, ERROR_CODE
```

`MPI_FILE_SET_ERRHANDLER(file, errhandler)`

INOUT	file	file (handle)
-------	------	---------------

IN	errhandler	new error handler for file (handle)
----	------------	-------------------------------------

### C binding

```
int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)
```

### Fortran 2008 binding

```
MPI_File_set_errhandler(file, errhandler, ierror)
```

```
  TYPE(MPI_File), INTENT(IN) :: file
```

```
  TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
```

```
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
```

```
  INTEGER FILE, ERRHANDLER, IERROR
```

Attaches a new error handler to a file. The error handler must be either a predefined error handler, or an error handler created by a call to `MPI_FILE_CREATE_ERRHANDLER`.

`MPI_FILE_GET_ERRHANDLER(file, errhandler)`

IN	file	file (handle)
----	------	---------------

OUT	errhandler	error handler currently associated with file (handle)
-----	------------	---

### C binding

```
int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)
```

**Fortran 2008 binding**

```

MPI_File_get_errhandler(file, errhandler, ierror)
    TYPE(MPI_File), INTENT(IN) :: file
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

**Fortran binding**

```

MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
    INTEGER FILE, ERRHANDLER, IERROR

```

Retrieves the error handler currently associated with a file.

**8.3.4 Error Handlers for Sessions**

```

MPI_SESSION_CREATE_ERRHANDLER(session_errhandler_fn, errhandler)

```

IN	session_errhandler_fn	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

**C binding**

```

int MPI_Session_create_errhandler(
    MPI_Session_errhandler_function *session_errhandler_fn,
    MPI_Errhandler *errhandler)

```

**Fortran 2008 binding**

```

MPI_Session_create_errhandler(session_errhandler_fn, errhandler, ierror)
    PROCEDURE(MPI_Session_errhandler_function) :: session_errhandler_fn
    TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

**Fortran binding**

```

MPI_SESSION_CREATE_ERRHANDLER(SESSION_ERRHANDLER_FN, ERRHANDLER, IERROR)
    EXTERNAL SESSION_ERRHANDLER_FN
    INTEGER ERRHANDLER, IERROR

```

Creates an error handler that can be attached to a session object. In C, the `session_errhandler_fn` argument should be a function of type `MPI_Session_errhandler_function`, which is defined as

```

typedef void MPI_Session_errhandler_function(MPI_Session *session,
    int *error_code, ...);

```

The first argument is the session in use, the second is the error code to be returned. The remaining arguments are “`varargs`” arguments whose number and meaning is implementation-dependent. An implementation should clearly document these arguments. With the Fortran `mpi_f08` module, the `session_errhandler_fn` argument should be of the form:

```

ABSTRACT INTERFACE
    SUBROUTINE MPI_Session_errhandler_function(session, error_code)
        TYPE(MPI_Session) :: session

```

```
INTEGER :: error_code
```

With the Fortran `mpi` module and `mpif.h`, the `SESSION_ERRHANDLER_FN` argument should be of the form:

```
SUBROUTINE SESSION_ERRHANDLER_FUNCTION(SESSION, ERROR_CODE)
  INTEGER SESSION, ERROR_CODE
```

```
MPI_SESSION_SET_ERRHANDLER(session, errhandler)
```

INOUT	session	session (handle)
IN	errhandler	new error handler for session (handle)

### C binding

```
int MPI_Session_set_errhandler(MPI_Session session,
                               MPI_Errhandler errhandler)
```

### Fortran 2008 binding

```
MPI_Session_set_errhandler(session, errhandler, ierror)
  TYPE(MPI_Session), INTENT(IN) :: session
  TYPE(MPI_Errhandler), INTENT(IN) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_SESSION_SET_ERRHANDLER(SESSION, ERRHANDLER, IERROR)
  INTEGER SESSION, ERRHANDLER, IERROR
```

Attaches a new error handler to a session. The error handler must be either a pre-defined error handler, or an error handler created by a call to `MPI_SESSION_CREATE_ERRHANDLER`.

```
MPI_SESSION_GET_ERRHANDLER(session, errhandler)
```

IN	session	session (handle)
OUT	errhandler	error handler currently associated with session (handle)

### C binding

```
int MPI_Session_get_errhandler(MPI_Session session,
                               MPI_Errhandler *errhandler)
```

### Fortran 2008 binding

```
MPI_Session_get_errhandler(session, errhandler, ierror)
  TYPE(MPI_Session), INTENT(IN) :: session
  TYPE(MPI_Errhandler), INTENT(OUT) :: errhandler
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_SESSION_GET_ERRHANDLER(SESSION, ERRHANDLER, IERROR)
  INTEGER SESSION, ERRHANDLER, IERROR
```

Retrieves the error handler currently associated with a session.

### 8.3.5 Freeing Errorhandlers and Retrieving Error Strings

**MPI\_ERRHANDLER\_FREE**(errhandler)

INOUT      errhandler                      MPI error handler (handle)

#### C binding

```
int MPI_Errhandler_free(MPI_Errhandler *errhandler)
```

#### Fortran 2008 binding

```
MPI_Errhandler_free(errhandler, ierror)
    TYPE(MPI_Errhandler), INTENT(INOUT) :: errhandler
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

#### Fortran binding

```
MPI_ERRHANDLER_FREE(ERRHANDLER, IERROR)
    INTEGER ERRHANDLER, IERROR
```

Marks the error handler associated with `errhandler` for deallocation and sets `errhandler` to `MPI_ERRHANDLER_NULL`. The error handler will be deallocated after all the objects associated with it (communicator, window, or file) have been deallocated.

**MPI\_ERROR\_STRING**(errorcode, string, resultlen)

IN	errorcode	Error code returned by an MPI routine
OUT	string	Text that corresponds to the errorcode
OUT	resultlen	Length (in printable characters) of the result returned in string

#### C binding

```
int MPI_Error_string(int errorcode, char *string, int *resultlen)
```

#### Fortran 2008 binding

```
MPI_Error_string(errorcode, string, resultlen, ierror)
    INTEGER, INTENT(IN) :: errorcode
    CHARACTER(LEN=MPI_MAX_ERROR_STRING), INTENT(OUT) :: string
    INTEGER, INTENT(OUT) :: resultlen
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

#### Fortran binding

```
MPI_ERROR_STRING(ERRORCODE, STRING, RESULTLEN, IERROR)
    INTEGER ERRORCODE, RESULTLEN, IERROR
    CHARACTER*(*) STRING
```

Returns the error string associated with an error code or class. The argument `string` must represent storage that is at least `MPI_MAX_ERROR_STRING` characters long.

The number of characters actually written is returned in the output argument, `resultlen`.

This function must always be thread-safe, as defined in Section 11.6. It is one of the few routines that may be called before MPI is initialized or after MPI is finalized.

*Rationale.* The form of this function was chosen to make the Fortran and C bindings similar. A version that returns a pointer to a string has two difficulties. First, the return string must be statically allocated and different for each error message (allowing the pointers returned by successive calls to `MPI_ERROR_STRING` to point to the correct message). Second, in Fortran, a function declared as returning `CHARACTER*(*)` can not be referenced in, for example, a `PRINT` statement. (*End of rationale.*)

## 8.4 Error Codes and Classes

The error codes returned by MPI are left entirely to the implementation (with the exception of `MPI_SUCCESS`). This is done to allow an implementation to provide as much information as possible in the error code (for use with `MPI_ERROR_STRING`).

All MPI function calls shall return `MPI_SUCCESS` if and only if the specification of that function has been fulfilled at the point of return. For multiple completion functions, if the function returns `MPI_ERR_IN_STATUS`, the error code in each status object shall be set to `MPI_SUCCESS` if and only if the specification of the operation represented by the corresponding `MPI_Request` has been fulfilled at the point of return.

When an operation raises an error, it may not satisfy its specification (for example, a synchronizing operation may not have synchronized) and the content of the output buffers, targeted memory, or output parameters is undefined. However, a valid error code shall always be set when an operation raises an error, whether in the return value, error field in the status object, or element in an array of error codes.

To make it possible for an application to interpret an error code, the routine `MPI_ERROR_CLASS` converts any error code into one of a small set of standard error codes, called *error classes*. Valid error classes are shown in Table 8.1 and Table 8.2.

The error classes are a subset of the error codes: an MPI function may return an error class number; and the function `MPI_ERROR_STRING` can be used to compute the error string associated with an error class. The values defined for MPI error classes are valid MPI error codes.

The error codes satisfy,

$$0 = \text{MPI\_SUCCESS} < \text{MPI\_ERR\_}\dots \leq \text{MPI\_ERR\_LASTCODE}.$$

*Rationale.* The difference between `MPI_ERR_UNKNOWN` and `MPI_ERR_OTHER` is that `MPI_ERROR_STRING` can return useful information about `MPI_ERR_OTHER`.

Note that `MPI_SUCCESS = 0` is necessary to be consistent with C practice; the separation of error classes and error codes allows us to define the error classes this way. Having a known `LASTCODE` is often a nice sanity check as well. (*End of rationale.*)

1	MPI_SUCCESS	No error
2	MPI_ERR_ACCESS	Permission denied
3	MPI_ERR_AMODE	Error related to the amode passed to
4		MPI_FILE_OPEN
5	MPI_ERR_ARG	Invalid argument of some other kind
6	MPI_ERR_ASSERT	Invalid assertion argument
7	MPI_ERR_BAD_FILE	Invalid file name (e.g., path name too long)
8	MPI_ERR_BASE	Invalid base passed to MPI_FREE_MEM
9	MPI_ERR_BUFFER	Invalid buffer pointer argument
10	MPI_ERR_COMM	Invalid communicator argument
11	MPI_ERR_CONVERSION	An error occurred in a user supplied data
12		conversion function
13	MPI_ERR_COUNT	Invalid count argument
14	MPI_ERR_DIMS	Invalid dimension argument
15	MPI_ERR_DISP	Invalid displacement argument
16	MPI_ERR_DUP_DATAREP	Conversion functions could not be regis-
17		tered because a data representation identi-
18		fier that was already defined was passed to
19		MPI_REGISTER_DATAREP
20	MPI_ERR_FILE	Invalid file handle argument
21	MPI_ERR_FILE_EXISTS	File exists
22	MPI_ERR_FILE_IN_USE	File operation could not be completed, as
23		the file is currently open by some process
24	MPI_ERR_GROUP	Invalid group argument
25	MPI_ERR_INFO	Invalid info argument
26	MPI_ERR_INFO_KEY	Key longer than MPI_MAX_INFO_KEY
27	MPI_ERR_INFO_NOKEY	Invalid key passed to MPI_INFO_DELETE
28	MPI_ERR_INFO_VALUE	Value longer than MPI_MAX_INFO_VAL
29	MPI_ERR_IN_STATUS	Error code is in status
30	MPI_ERR_INTERN	Internal MPI (implementation) error
31	MPI_ERR_IO	Other I/O error
32	MPI_ERR_KEYVAL	Invalid keyval argument
33	MPI_ERR_LOCKTYPE	Invalid locktype argument
34	MPI_ERR_NAME	Invalid service name passed to
35		MPI_LOOKUP_NAME
36	MPI_ERR_NO_MEM	MPI_ALLOC_MEM failed because memory
37		is exhausted
38	MPI_ERR_NO_SPACE	Not enough space
39	MPI_ERR_NO_SUCH_FILE	File does not exist
40	MPI_ERR_NOT_SAME	Collective argument not identical on all
41		processes, or collective routines called in
42		a different order by different processes
43		
44		
45		
46		
47		
48		

Table 8.1: Error classes (Part 1)



MPI_ERR_OP	Invalid operation argument	1
MPI_ERR_OTHER	Known error not in this list	2
MPI_ERR_PENDING	Pending request	3
MPI_ERR_PORT	Invalid port name passed to MPI_COMM_CONNECT	4 5
MPI_ERR_PROC_ABORTED	Operation failed because a peer process has aborted	6 7
MPI_ERR_QUOTA	Quota exceeded	8
MPI_ERR_RANK	Invalid rank argument	9
MPI_ERR_READ_ONLY	Read-only file or file system	10
MPI_ERR_REQUEST	Invalid request argument	11
MPI_ERR_RMA_ATTACH	Memory cannot be attached (e.g., because of resource exhaustion)	12 13
MPI_ERR_RMA_CONFLICT	Conflicting accesses to window	14
MPI_ERR_RMA_FLAVOR	Passed window has the wrong flavor for the called function	15 16
MPI_ERR_RMA_RANGE	Target memory is not part of the win- dow (in the case of a window created with MPI_WIN_CREATE_DYNAMIC, tar- get memory is not attached)	17 18 19 20
MPI_ERR_RMA_SHARED	Memory cannot be shared (e.g., some pro- cess in the group of the specified commu- nicator cannot expose shared memory)	21 22 23
MPI_ERR_RMA_SYNC	Wrong synchronization of RMA calls	24
MPI_ERR_ROOT	Invalid root argument	25
MPI_ERR_SERVICE	Invalid service name passed to MPI_UNPUBLISH_NAME	26 27
MPI_ERR_SESSION	Invalid session argument	28
MPI_ERR_SIZE	Invalid size argument	29
MPI_ERR_SPAWN	Error in spawning processes	30
MPI_ERR_TAG	Invalid tag argument	31
MPI_ERR_TOPOLOGY	Invalid topology argument	32
MPI_ERR_TRUNCATE	Message truncated on receive	33
MPI_ERR_TYPE	Invalid datatype argument	34
MPI_ERR_UNKNOWN	Unknown error	35
MPI_ERR_UNSUPPORTED_DATAREP	Unsupported datarep passed to MPI_FILE_SET_VIEW	36 37
MPI_ERR_UNSUPPORTED_OPERATION	Unsupported operation, such as seeking on a file which supports sequential access only	38 39
MPI_ERR_VALUE_TOO_LARGE	Value is too large to store	40
MPI_ERR_WIN	Invalid window argument	41
MPI_ERR_LASTCODE	Last error code	42

Table 8.2: Error classes (Part 2)

```
1 MPI_ERROR_CLASS(errorcode, errorclass)
```

```
2     IN          errorcode          Error code returned by an MPI routine
```

```
3     OUT        errorclass         Error class associated with errorcode
```

### 6 C binding

```
7 int MPI_Error_class(int errorcode, int *errorclass)
```

### 9 Fortran 2008 binding

```
10 MPI_Error_class(errorcode, errorclass, ierror)
```

```
11     INTEGER, INTENT(IN) :: errorcode
```

```
12     INTEGER, INTENT(OUT) :: errorclass
```

```
13     INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### 14 Fortran binding

```
15 MPI_ERROR_CLASS(ERRORCODE, ERRORCLASS, IERROR)
```

```
16     INTEGER ERRORCODE, ERRORCLASS, IERROR
```

17 The function `MPI_ERROR_CLASS` maps each standard error code (error class) onto  
18 itself.

19 This function must always be thread-safe, as defined in Section 11.6. It is one of the  
20 few routines that may be called before MPI is initialized or after MPI is finalized.  
21

## 23 8.5 Error Classes, Error Codes, and Error Handlers

24 Users may want to write a layered library on top of an existing MPI implementation, and  
25 this library may have its own set of error codes and classes. An example of such a library  
26 is an I/O library based on MPI, see Chapter 14. For this purpose, functions are needed to:  
27

- 28 1. add a new error class to the ones an MPI implementation already knows.
- 29 2. associate error codes with this error class, so that `MPI_ERROR_CLASS` works.
- 30 3. associate strings with these error codes, so that `MPI_ERROR_STRING` works.
- 31 4. invoke the error handler associated with a communicator, window, or object.
- 32
- 33
- 34

35 Several functions are provided to do this. They are all local. No functions are provided  
36 to free error classes or codes: it is not expected that an application will generate them in  
37 significant numbers.  
38

```
39 MPI_ADD_ERROR_CLASS(errorclass)
```

```
40     OUT        errorclass         value for the new error class (integer)
```

### 43 C binding

```
44 int MPI_Add_error_class(int *errorclass)
```

### 46 Fortran 2008 binding

```
47 MPI_Add_error_class(errorclass, ierror)
```

```
48     INTEGER, INTENT(OUT) :: errorclass
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)
```

```
INTEGER ERRORCLASS, IERROR
```

Creates a new error class and returns the value for it.

*Rationale.* To avoid conflicts with existing error codes and classes, the value is set by the implementation and not by the user. (*End of rationale.*)

*Advice to users.* Since a call to `MPI_ADD_ERROR_CLASS` is local, the same `errorclass` may not be returned on all processes that make this call. Thus, it is not safe to assume that registering a new error on a set of processes at the same time will yield the same `errorclass` on all of the processes. Getting the “same” error on multiple processes may not cause the same value of error code to be generated. (*End of advice to users.*)

The value of `MPI_ERR_LASTCODE` is a constant value and is not affected by new user-defined error codes and classes. Instead, a predefined attribute key `MPI_LASTUSED` is associated with `MPI_COMM_WORLD`. The attribute value corresponding to this key is the current maximum error class including the user-defined ones. This is a local value and may be different on different processes. The value returned by this key is always greater than or equal to `MPI_ERR_LASTCODE`.

*Advice to users.* The value returned by the key `MPI_LASTUSED` will not change unless the user calls a function to explicitly add an error class/code. In a multithreaded environment, the user must take extra care in assuming this value has not changed. Note that error codes and error classes are not necessarily dense. A user may not assume that each error class below `MPI_LASTUSED` is valid. (*End of advice to users.*)

```
MPI_ADD_ERROR_CODE(errorclass, errorcode)
```

IN	errorclass	error class (integer)
OUT	errorcode	new error code to be associated with errorclass (integer)

### C binding

```
int MPI_Add_error_code(int errorclass, int *errorcode)
```

### Fortran 2008 binding

```
MPI_Add_error_code(errorclass, errorcode, ierror)
```

```
INTEGER, INTENT(IN) :: errorclass
```

```
INTEGER, INTENT(OUT) :: errorcode
```

```
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

### Fortran binding

```
MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)
```

```
INTEGER ERRORCLASS, ERRORCODE, IERROR
```

Creates new error code associated with `errorclass` and returns its value in `errorcode`.

*Rationale.* To avoid conflicts with existing error codes and classes, the value of the new error code is set by the implementation and not by the user. (*End of rationale.*)

**MPI\_ADD\_ERROR\_STRING**(errorcode, string)

IN	errorcode	error code or class (integer)
IN	string	text corresponding to errorcode (string)

### C binding

int MPI\_Add\_error\_string(int errorcode, const char \*string)

### Fortran 2008 binding

MPI\_Add\_error\_string(errorcode, string, ierror)

```

INTEGER, INTENT(IN) :: errorcode
CHARACTER(LEN=*), INTENT(IN) :: string
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

### Fortran binding

MPI\_ADD\_ERROR\_STRING(ERRORCODE, STRING, IERROR)

```

INTEGER ERRORCODE, IERROR
CHARACTER*(*) STRING

```

Associates an error string with an error code or class. The string must be no more than MPI\_MAX\_ERROR\_STRING characters long. The length of the string is as defined in the calling language. The length of the string does not include the null terminator in C. Trailing blanks will be stripped in Fortran. Calling MPI\_ADD\_ERROR\_STRING for an errorcode that already has a string will replace the old string with the new string. It is erroneous to call MPI\_ADD\_ERROR\_STRING for an error code or class with a value  $\leq$  MPI\_ERR\_LASTCODE.

If MPI\_ERROR\_STRING is called when no string has been set, it will return a empty string (all spaces in Fortran, "" in C).

Section 8.3 describes the methods for creating and associating error handlers with communicators, files, windows, and sessions.

**MPI\_COMM\_CALL\_ERRHANDLER**(comm, errorcode)

IN	comm	communicator with error handler (handle)
IN	errorcode	error code (integer)

### C binding

int MPI\_Comm\_call\_errhandler(MPI\_Comm comm, int errorcode)

### Fortran 2008 binding

MPI\_Comm\_call\_errhandler(comm, errorcode, ierror)

```

TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, INTENT(IN) :: errorcode
INTEGER, OPTIONAL, INTENT(OUT) :: ierror

```

**Fortran binding**

```
MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)
    INTEGER COMM, ERRORCODE, IERROR
```

This function invokes the error handler assigned to the communicator with the error code supplied. This function returns `MPI_SUCCESS` in C and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

```
MPI_WIN_CALL_ERRHANDLER(win, errorcode)
```

```
IN      win                window with error handler (handle)
IN      errorcode          error code (integer)
```

**C binding**

```
int MPI_Win_call_errhandler(MPI_Win win, int errorcode)
```

**Fortran 2008 binding**

```
MPI_Win_call_errhandler(win, errorcode, ierror)
    TYPE(MPI_Win), INTENT(IN) :: win
    INTEGER, INTENT(IN) :: errorcode
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

**Fortran binding**

```
MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)
    INTEGER WIN, ERRORCODE, IERROR
```

This function invokes the error handler assigned to the window with the error code supplied. This function returns `MPI_SUCCESS` in C and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

*Advice to users.* In contrast to communicators, the error handler `MPI_ERRORS_ARE_FATAL` is associated with a window when it is created. (*End of advice to users.*)

```
MPI_FILE_CALL_ERRHANDLER(fh, errorcode)
```

```
IN      fh                file with error handler (handle)
IN      errorcode          error code (integer)
```

**C binding**

```
int MPI_File_call_errhandler(MPI_File fh, int errorcode)
```

**Fortran 2008 binding**

```
MPI_File_call_errhandler(fh, errorcode, ierror)
    TYPE(MPI_File), INTENT(IN) :: fh
    INTEGER, INTENT(IN) :: errorcode
```

```
1      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

## Fortran binding

```
3      MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)
4      INTEGER FH, ERRORCODE, IERROR
```

6 This function invokes the error handler assigned to the file with the error code supplied.  
7 This function returns MPI\_SUCCESS in C and the same value in IERROR if the error handler  
8 was successfully called (assuming the process is not aborted and the error handler returns).

10 *Advice to users.* The default error handler for files is MPI\_ERRORS\_RETURN. (*End of*  
11 *advice to users.*)

```
14 MPI_SESSION_CALL_ERRHANDLER(session, errorcode)
```

```
16      IN          session          session with error handler (handle)
17      IN          errorcode        error code (integer)
```

## C binding

```
20 int MPI_Session_call_errhandler(MPI_Session session, int errorcode)
```

## Fortran 2008 binding

```
23 MPI_Session_call_errhandler(session, errorcode, ierror)
24      TYPE(MPI_Session), INTENT(IN) :: session
25      INTEGER, INTENT(IN) :: errorcode
26      INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

## Fortran binding

```
28 MPI_SESSION_CALL_ERRHANDLER(SESSION, ERRORCODE, IERROR)
29      INTEGER SESSION, ERRORCODE, IERROR
```

31 This function invokes the error handler assigned to the session with the error code  
32 supplied. This function returns MPI\_SUCCESS in C and the same value in IERROR if the  
33 error handler was successfully called (assuming the process is not aborted and the error  
34 handler returns).

36 *Advice to users.* Users are warned that handlers should not be called recursively  
37 with MPI\_COMM\_CALL\_ERRHANDLER, MPI\_FILE\_CALL\_ERRHANDLER,  
38 MPI\_WIN\_CALL\_ERRHANDLER, or MPI\_SESSION\_CALL\_ERRHANDLER. Doing this  
39 can create a situation where an infinite recursion is created. This can occur if  
40 MPI\_COMM\_CALL\_ERRHANDLER, MPI\_FILE\_CALL\_ERRHANDLER,  
41 MPI\_WIN\_CALL\_ERRHANDLER, or MPI\_SESSION\_CALL\_ERRHANDLER is called in-  
42 side an error handler.

43 Error codes and classes are associated with a process. As a result, they may be used  
44 in any error handler. Error handlers should be prepared to deal with any error code  
45 they are given. Furthermore, it is good practice to only call an error handler with the  
46 appropriate error codes. For example, file errors would normally be sent to the file  
47 error handler. (*End of advice to users.*)

## 8.6 Timers and Synchronization

MPI defines a timer. A timer is specified even though it is not “message-passing,” because timing parallel programs is important in “performance debugging” and because existing timers (both in POSIX 1003.1-1988 and 1003.4D 14.1 and in Fortran 90) are either inconvenient or do not provide adequate access to high resolution timers. See also Section 2.6.4.

**MPI\_WTIME()**

**C binding**

```
double MPI_Wtime(void)
```

**Fortran 2008 binding**

```
DOUBLE PRECISION MPI_Wtime()
```

**Fortran binding**

```
DOUBLE PRECISION MPI_WTIME()
```

**MPI\_WTIME** returns a floating-point number of seconds, representing elapsed wall-clock time since some time in the past.

The “time in the past” is guaranteed not to change during the life of the process. The user is responsible for converting large numbers of seconds to other units if they are preferred.

This function is portable (it returns seconds, not “ticks”), and it allows high-resolution. One would use it like this:

```
{
    double starttime, endtime;
    starttime = MPI_Wtime();
    ... stuff to be timed ...
    endtime   = MPI_Wtime();
    printf("That took %f seconds\n", endtime-starttime);
}
```

The times returned are local to the node that called them. There is no requirement that different nodes return “the same time.” (But see also the discussion of **MPI\_WTIME\_IS\_GLOBAL** in Section 8.1.2).

**MPI\_WTICK()**

**C binding**

```
double MPI_Wtick(void)
```

**Fortran 2008 binding**

```
DOUBLE PRECISION MPI_Wtick()
```

**Fortran binding**

```
DOUBLE PRECISION MPI_WTICK()
```

MPI\_WTICK returns the resolution of MPI\_WTIME in seconds. That is, it returns, as a double precision value, the number of seconds between successive clock ticks. For example, if the clock is implemented by the hardware as a counter that is incremented every millisecond, the value returned by MPI\_WTICK should be  $(10^{-3})$ .



# Index

MPI\_ANY\_SOURCE, [4](#)  
MPI\_COMM\_SELF, [10](#)  
MPI\_COMM\_TYPE\_HW\_GUIDED, [6](#), [7](#)  
MPI\_COMM\_WORLD, [3](#), [4](#), [13](#), [25](#)  
MPI\_ERR\_ACCESS, [22](#)  
MPI\_ERR\_AMODE, [22](#)  
MPI\_ERR\_ARG, [22](#)  
MPI\_ERR\_ASSERT, [22](#)  
MPI\_ERR\_BAD\_FILE, [22](#)  
MPI\_ERR\_BASE, [9](#), [22](#)  
MPI\_ERR\_BUFFER, [22](#)  
MPI\_ERR\_COMM, [22](#)  
MPI\_ERR\_CONVERSION, [22](#)  
MPI\_ERR\_COUNT, [22](#)  
MPI\_ERR\_DIMS, [22](#)  
MPI\_ERR\_DISP, [22](#)  
MPI\_ERR\_DUP\_DATAREP, [22](#)  
MPI\_ERR\_FILE, [22](#)  
MPI\_ERR\_FILE\_EXISTS, [22](#)  
MPI\_ERR\_FILE\_IN\_USE, [22](#)  
MPI\_ERR\_GROUP, [22](#)  
MPI\_ERR\_IN\_STATUS, [13](#), [21](#), [22](#)  
MPI\_ERR\_INFO, [22](#)  
MPI\_ERR\_INFO\_KEY, [22](#)  
MPI\_ERR\_INFO\_NOKEY, [22](#)  
MPI\_ERR\_INFO\_VALUE, [22](#)  
MPI\_ERR\_INTERN, [11](#), [22](#)  
MPI\_ERR\_IO, [22](#)  
MPI\_ERR\_KEYVAL, [22](#)  
MPI\_ERR\_LASTCODE, [21](#), [23](#), [25](#), [25](#), [26](#)  
MPI\_ERR\_LOCKTYPE, [22](#)  
MPI\_ERR\_NAME, [22](#)  
MPI\_ERR\_NO\_MEM, [8](#), [22](#)  
MPI\_ERR\_NO\_SPACE, [22](#)  
MPI\_ERR\_NO\_SUCH\_FILE, [22](#)  
MPI\_ERR\_NOT\_SAME, [22](#)  
MPI\_ERR\_OP, [23](#)  
MPI\_ERR\_OTHER, [21](#), [23](#)  
MPI\_ERR\_PENDING, [23](#)  
MPI\_ERR\_PORT, [23](#)  
MPI\_ERR\_PROC\_ABORTED, [23](#)  
MPI\_ERR\_QUOTA, [23](#)  
MPI\_ERR\_RANK, [23](#)  
MPI\_ERR\_READ\_ONLY, [23](#)  
MPI\_ERR\_REQUEST, [23](#)  
MPI\_ERR\_RMA\_ATTACH, [23](#)  
MPI\_ERR\_RMA\_CONFLICT, [23](#)  
MPI\_ERR\_RMA\_FLAVOR, [23](#)  
MPI\_ERR\_RMA\_RANGE, [23](#)  
MPI\_ERR\_RMA\_SHARED, [23](#)  
MPI\_ERR\_RMA\_SYNC, [23](#)  
MPI\_ERR\_ROOT, [23](#)  
MPI\_ERR\_SERVICE, [23](#)  
MPI\_ERR\_SESSION, [23](#)  
MPI\_ERR\_SIZE, [23](#)  
MPI\_ERR\_SPAWN, [23](#)  
MPI\_ERR\_TAG, [23](#)  
MPI\_ERR\_TOPOLOGY, [23](#)  
MPI\_ERR\_TRUNCATE, [23](#)  
MPI\_ERR\_TYPE, [23](#)  
MPI\_ERR\_UNKNOWN, [21](#), [23](#)  
MPI\_ERR\_UNSUPPORTED\_DATAREP, [23](#)  
MPI\_ERR\_UNSUPPORTED\_OPERATION, [23](#)  
MPI\_ERR\_VALUE\_TOO\_LARGE, [23](#)  
MPI\_ERR\_WIN, [23](#)  
MPI\_Errhandler, [12](#), [13–20](#)  
MPI\_ERRHANDLER\_NULL, [20](#)  
MPI\_ERRORS\_ABORT, [11](#), [11](#)  
MPI\_ERRORS\_ARE\_FATAL, [11](#), [11](#), [12](#), [27](#)  
MPI\_ERRORS\_RETURN, [11](#), [11](#), [12](#), [28](#)  
MPI\_File, [17](#)  
MPI\_Get\_hw\_resource\_status, [6](#)  
MPI\_Get\_hw\_resource\_types, [5](#)  
MPI\_HOST, [3](#), [3](#)  
MPI\_HW\_OCCUPIED, [7](#)  
MPI\_HW\_PRESENT, [7](#)  
MPI\_HW\_UNKNOWN, [7](#)  
MPI\_HW\_USABLE, [7](#)  
MPI\_Info, [7](#)  
MPI\_INFO\_NULL, [8](#)

- 1 MPI\_IO, [3](#), [4](#)
- 2 MPI\_LASTUSED\_CODE, [25](#)
- 3 MPI\_MAX\_ERROR\_STRING, [20](#), [26](#)
- 4 MPI\_MAX\_INFO\_KEY, [22](#)
- 5 MPI\_MAX\_INFO\_VAL, [22](#)
- 6 MPI\_MAX\_LIBRARY\_VERSION\_STRING, [2](#), [2](#)
- 7 MPI\_MAX\_PROCESSOR\_NAME, [5](#), [5](#)
- 8 MPI\_PROC\_NULL, [3](#), [4](#)
- 9 MPI\_Request, [21](#)
- 10 MPI\_Session, [18](#), [19](#)
- 11 MPI\_SUBVERSION, [2](#)
- 12 MPI\_SUCCESS, [21](#), [22](#), [27](#), [28](#)
- 13 MPI\_TAG\_UB, [3](#), [3](#)
- 14 MPI\_VERSION, [2](#)
- 15 MPI\_Win, [15](#), [16](#)
- 16 MPI\_WTIME\_IS\_GLOBAL, [3](#), [4](#), [29](#)
- 17 TYPE(MPI\_Errhandler), [12](#)
- 18
- 19 EXAMPLES:MPI\_ALLOC\_MEM, [9](#)
- 20 EXAMPLES:MPI\_Alloc\_mem, [10](#)
- 21 EXAMPLES:MPI\_FREE\_MEM, [9](#)
- 22
- 23 "mpi\_hw\_res\_i\_alias\_k", [6](#)
- 24 "mpi\_hw\_res\_i\_aliases", [6](#)
- 25 "mpi\_hw\_res\_i\_occupied", [6](#)
- 26 "mpi\_hw\_res\_i\_type", [6](#)
- 27 "mpi\_hw\_res\_j\_type", [6](#)
- 28 "mpi\_hw\_res\_i\_aliases", [6](#)
- 29 "mpi\_hw\_res\_nresources", [6](#)
- 30 "mpi\_hw\_resource\_type", [6](#), [7](#)
- 31 "mpi\_minimum\_memory\_alignment", [8](#)
- 32 "false", [6](#)
- 33 "true", [6](#)
- 34 MPI\_ABORT, [11](#)
- 35 MPI\_ADD\_ERROR\_CLASS, [25](#)
- 36 x errorclass), [24](#)
- 37 MPI\_ADD\_ERROR\_STRING, [26](#)
- 38 MPI\_ALLOC\_MEM, [8](#), [9](#), [22](#)
- 39 MPI\_ALLOC\_MEM\_CPTR, [8](#)
- 40 MPI\_COMM\_CALL\_ERRHANDLER, [28](#)
- 41 MPI\_COMM\_CONNECT, [23](#)
- 42 MPI\_COMM\_CREATE\_ERRHANDLER, [12](#), [14](#)
- 43
- 44 MPI\_COMM\_CREATE\_FROM\_GROUP, [3](#)
- 45 MPI\_COMM\_GET\_ATTR, [3](#)
- 46 MPI\_COMM\_GET\_ERRHANDLER, [12](#)
- 47 MPI\_COMM\_GROUP, [12](#)
- 48 MPI\_COMM\_SET\_ERRHANDLER, [12](#)
- MPI\_COMM\_SPLIT\_TYPE, [6](#), [7](#)
- MPI\_ERRHANDLER\_FREE, [12](#)
- x errhandler), [20](#)
- MPI\_ERROR\_CLASS, [21](#), [24](#)
- MPI\_ERROR\_STRING, [21](#), [24](#), [26](#)
- MPI\_FILE\_CALL\_ERRHANDLER, [28](#)
- MPI\_FILE\_CREATE\_ERRHANDLER, [12](#), [17](#)
- MPI\_FILE\_GET\_ERRHANDLER, [12](#)
- MPI\_FILE\_OPEN, [22](#)
- MPI\_FILE\_SET\_ERRHANDLER, [12](#)
- MPI\_FILE\_SET\_VIEW, [23](#)
- MPI\_FINALIZE, [3](#), [10](#)
- MPI\_FREE\_MEM, [9](#), [22](#)
- x base), [8](#)
- MPI\_GET\_HW\_RESOURCE\_TYPES, [5](#), [6](#)
- x hw\_info), [5](#)
- MPI\_GET\_LIBRARY\_VERSION, [2](#)
- MPI\_GET\_PROCESSOR\_NAME, [5](#)
- MPI\_GET\_VERSION, [2](#)
- MPI\_GROUP\_FREE, [12](#)
- MPI\_INFO\_DELETE, [22](#)
- MPI\_INFO\_FREE, [6](#)
- MPI\_INIT, [3](#), [10](#)
- MPI\_INIT\_THREAD, [10](#)
- MPI\_INTERCOMM\_CREATE\_FROM\_GROUPS, [3](#)
- MPI\_LOOKUP\_NAME, [22](#)
- MPI\_REGISTER\_DATAREP, [22](#)
- MPI\_SESSION\_CALL\_ERRHANDLER, [28](#)
- MPI\_SESSION\_CREATE\_ERRHANDLER, [12](#), [19](#)
- MPI\_SESSION\_GET\_ERRHANDLER, [12](#)
- MPI\_SESSION\_INIT, [12](#)
- MPI\_SESSION\_SET\_ERRHANDLER, [12](#)
- MPI\_UNPUBLISH\_NAME, [23](#)
- MPI\_WIN\_CALL\_ERRHANDLER, [28](#)
- MPI\_WIN\_CREATE\_DYNAMIC, [23](#)
- MPI\_WIN\_CREATE\_ERRHANDLER, [12](#), [16](#)
- MPI\_WIN\_GET\_ERRHANDLER, [12](#)
- MPI\_WIN\_SET\_ERRHANDLER, [12](#)
- MPI\_WTICK, [30](#)
- MPI\_WTICK(), [29](#)
- MPI\_WTIME, [4](#), [29](#), [30](#)
- MPI\_WTIME(), [29](#)
- TERM:alignment, [8](#)
- TERM:clock synchronization, [4](#)

TERM:environmental inquiries, <a href="#">3</a>	1
TERM:error handling, <a href="#">10</a>	2
error codes and classes, <a href="#">21</a> , <a href="#">24</a>	3
error handlers, <a href="#">12</a> , <a href="#">24</a>	4
initial error handler, <a href="#">10</a> , <a href="#">11</a>	5
TERM:host rank, <a href="#">3</a>	6
TERM:IO rank, <a href="#">4</a>	7
TERM:memory	8
alignment, <a href="#">8</a>	9
allocation, <a href="#">7</a>	10
TERM:MPI process initialization	11
Sessions Model, <a href="#">3</a>	12
World Model, <a href="#">3</a>	13
TERM:POSIX	14
environment, <a href="#">29</a>	15
TERM:processor name, <a href="#">4</a>	16
TERM:Sessions Model, <a href="#">3</a>	17
TERM:tag values, <a href="#">3</a>	18
TERM:threads	19
thread-safe, <a href="#">2</a> , <a href="#">21</a> , <a href="#">24</a>	20
TERM:timers and synchronization, <a href="#">29</a>	21
TERM:version inquiries, <a href="#">1</a>	22
TERM:World Model, <a href="#">3</a>	23
COMM_ERRHANDLER_FUNCTION, <a href="#">13</a> , <a href="#">13</a>	24
FILE_ERRHANDLER_FUNCTION, <a href="#">16</a> , <a href="#">17</a>	25
MPI_Comm_errhandler_function, <a href="#">12</a> , <a href="#">13</a>	26
MPI_Comm_errhandler_function( MPI_Comm *comm, int *error_code, ...), <a href="#">13</a>	27 28 29
MPI_File_errhandler_function, <a href="#">16</a> , <a href="#">17</a>	30
MPI_File_errhandler_function( MPI_File *file, int *error_code, ...), <a href="#">17</a>	31 32 33
MPI_Session_errhandler_function, <a href="#">18</a> , <a href="#">19</a>	34
MPI_Session_errhandler_function( MPI_Session *session, int *error_code, ...), <a href="#">18</a>	35 36 37
MPI_Win_errhandler_function, <a href="#">14</a> , <a href="#">15</a>	38
MPI_Win_errhandler_function( MPI_Win *win, int *error_code, ...), <a href="#">15</a>	39 40 41
SESSION_ERRHANDLER_FUNCTION, <a href="#">18</a> , <a href="#">19</a>	42
WIN_ERRHANDLER_FUNCTION, <a href="#">15</a> , <a href="#">15</a>	43 44 45 46 47 48