


Support for Complex Hardware Topologies in MPI

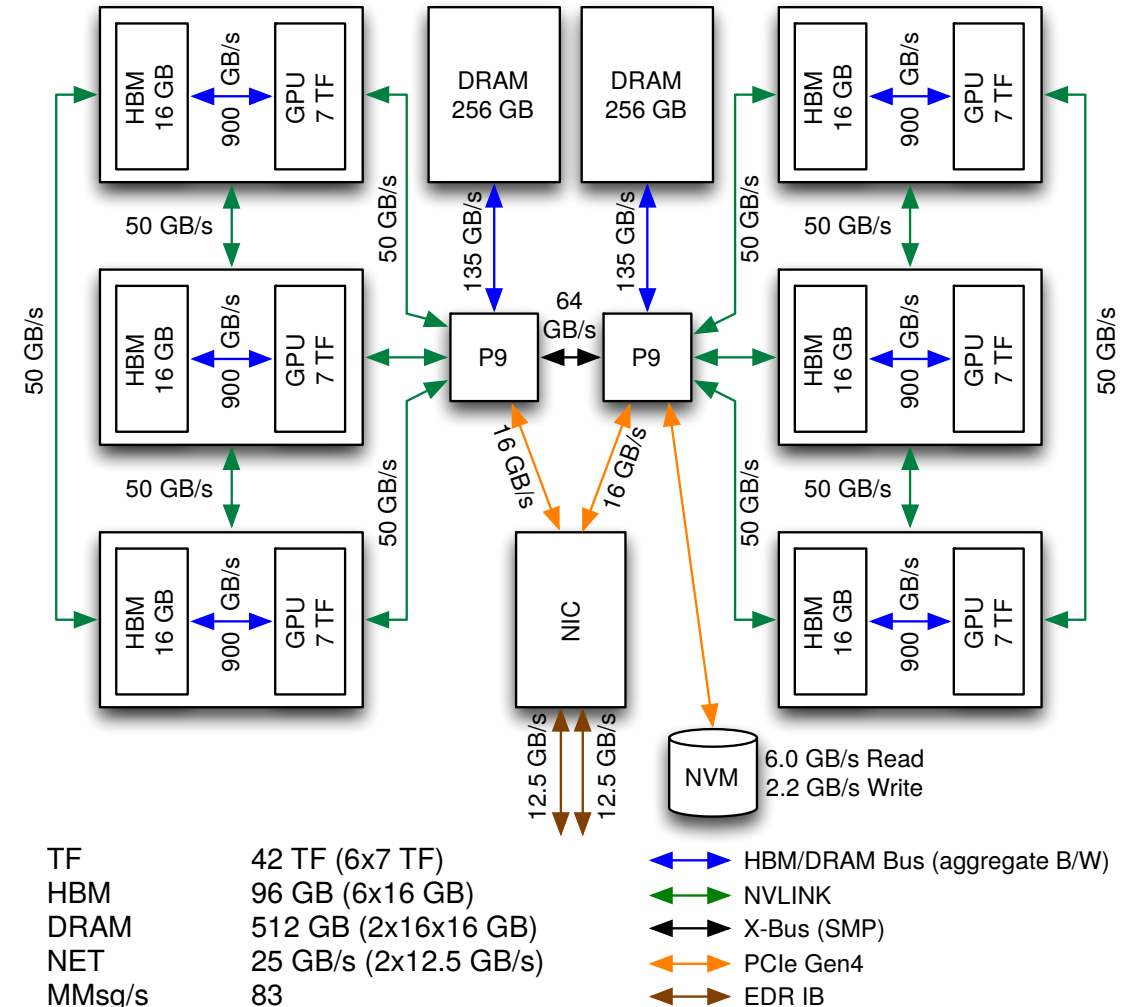
Geoffroy Vallee (ORNL)



Background

- Current and upcoming hardware are becoming more complex
 - Various hardware links: PCIe, NV Link, Gen-Z, etc.
 - Various memory technologies: NVe, HBM, DDR, etc.
- Hardware architectures now based on graph representations (not trees)
 - On Summit, NV Link is already creating a loop between GPUs
 - Cannot be represented as a tree anymore

How can we expose such architectures to applications through MPI?



HBM & DRAM speeds are aggregate (Read+Write).
All other speeds (X-Bus, NVLink, PCIe, IB) are bi-directional.

MPI – What already exists?

- MPI already supports using concepts:
 - Concept of topology
 - Concept of distributed graph
 - Groups of MPI processes
- Graph topologies
 - Provides a mean to organize MPI processes
 - Limitations
 - Virtual topology (*I tell MPI about the processes I know*)
 - Always create a communicator (can be expensive)
 - `MPI_DIST_GRAPH_CREATE(comm_old, n, sources, degrees, destinations, weights, info, reorder, comm_dist_graph)`

What would be beneficial?


- Get all MPI processes running “near by” – *MPI tell me about the processes that share resources with me*
 - On the same L1, L2, L3, NUMA node, ...etc.
 - No implicit creation of communicators which may never be used
 - Example
 - MPI rank is sharing L1 with rank 42
 - MPI rank is sharing NUMA nodes with ranks 42, 44, 45, 80
 - MPI rank is sharing the compute node with ranks 42, 44, 45, 80, 101, 110, 122
- Try to leverage concepts from other proposals
- Can support future architectures

Proposal 1 – Abstracting the underlying architecture


- Rely on groups
- Abstract the actually hardware architecture

Proposal 1 – Conceptual example

NICO



Numa Node 0



Numa Node 1

mynode

- We assume MPI processes are numbered sequentially starting from 0
- Green processes get at least 3 groups
 - 0, 1, 2, 3 with URI “uri://core_0”
 - 0, 1, 2, 3, 4 with URI “uri://numa_0”
 - 0, 1, 2, 3, 4, 5, 6, 7 with URI “uri://mynode”
- Blue process gets 2 groups
 - 0, 1, 2, 3, 4 with URI “uri://numa_0”
 - 0, 1, 2, 3, 4, 5, 6, 7 with URI “uri://mynode”
- Red processes get 2 groups
 - 5, 6, 7 with URI “uri://numa_1”
 - 0, 1, 2, 3, 4, 5, 6, 7 with URI “uri://mynode”

Proposal 1 – New MPI function

MPI_Hw_groups_get (MPI_Group groups[], uris, sizes)

Local operation that queries the runtime about group sharing resources

- OUT groups Array of groups representing MPI processes sharing hardware resources
- OUT uris Array of unique group identifiers, assigned by the runtime, representing the node type in hardware graph
- OUT sizes Number of groups

Note: uris are a similar concept than the one proposed in other proposals, but create a naming issue

Proposal 1 – Code example (1)

- Create communicators for all hardware levels involved in resource sharing

```
MPI_Hw_groups_get (&groups, &labels, &num_gps);  
MPI_Comm *new_comms = malloc (num_gps * sizeof (MPI_Comm));  
for (i = 0; i < num_gps; i++)  
{  
    MPI_Comm_create_from_group (groups[i], &new_comms[i]); /* Function from the session proposal */  
}
```


Proposal 1 – Code example (2)

- Create a communicator for all MPI processes on the same NUMA node

```
MPI_Hw_groups_get (&groups, &labels, &num_gps);
MPI_Comm *new_comms = malloc (num_gps * sizeof (MPI_Comm));
for (i = 0; i < num_gps; i++)
{
    if (uricmp (uris[i], "uri://compute_node") == 0)
        MPI_Comm_create_from_group (groups[i], &new_comm);
}
```

Proposal 1 – Benefits/Limitation Analysis

- Pros

- Guarantee separation between resource manager, runtime and the MPI standard
- Local operation / No creation of global handles/objects
- Leverage many concepts from the Session proposal

- Cons

- Node centric
- URI concept (from the Session proposal) is still unclear, lacking useful APIs
- The underlying graph is not directly exposed to applications (possible but potentially require more extensive new APIs)
- Does not expose the potential speed different for the link between MPI processes (PCIe versus NVLink)

Proposal 2 – Exposing the underlying architecture

- Rely on existing MPI functions
- Add a single new function that gets the distributed graph from group of MPI processes
- Require a single new function, "reverse" of MPI_Dist_graph_create()

`MPI_Group_get_hw_dist_graph (gp, in, src[], src_w[], out, dest[], dest_w[])`

Returns data that represents the underlying hardware through a distributed graph for a given group.

IN gp Group from which we want to the distributed graph representing the underlying hardware.

OUT in Size of src and src_w

OUT src MPI processes for which the calling process is a destination

OUT src_w Weights of the edges into the calling process (hardware relative distance)

OUT out Size of dest and dest_w

OUT dest MPI processes for which the calling process is a source

OUT dest_w Weights of the edges from the calling process (hardware relative distance)


Illustration

- Get local representation of all local MPI processes

```
/* Get the MPI processes that are running on the compute node */  
MPI_Comm_split_type (MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED, 0, NULL, &my_local_comm);  
  
/* Get group of local MPI processes */  
MPI_Comm_group (my_local_comm, &my_local_gp);  
  
/* Get the distributed graph from the group */  
MPI_Group_get_hw_dist_graph (my_local_gp, &in, &src[], &src_w[], &out, &dest[], &dest_w);
```

Higher levels in the hardware architecture

Lower levels in the hardware architecture



Proposal 2 – Benefits/Limitation Analysis

- Pros

- Guarantee separation between resource manager, runtime and the MPI standard
- Can be used to expose more than local resources (network neighbors)
- Expose the underlying hardware
- Scalable
- No naming issue (URIs)

- Cons

- MPI does not provide a type for graphs so it handles a graphs through arrays
- Like with MPI_Dist_graph, we only provide a distributed (but scalable) representation of the hardware